

Introduction to Data Science

Lecture 2


Data Collection and Exploration

Yesi N. Kunang



Outline

- Data Types and Sources
- Data Preparation
- Exploration



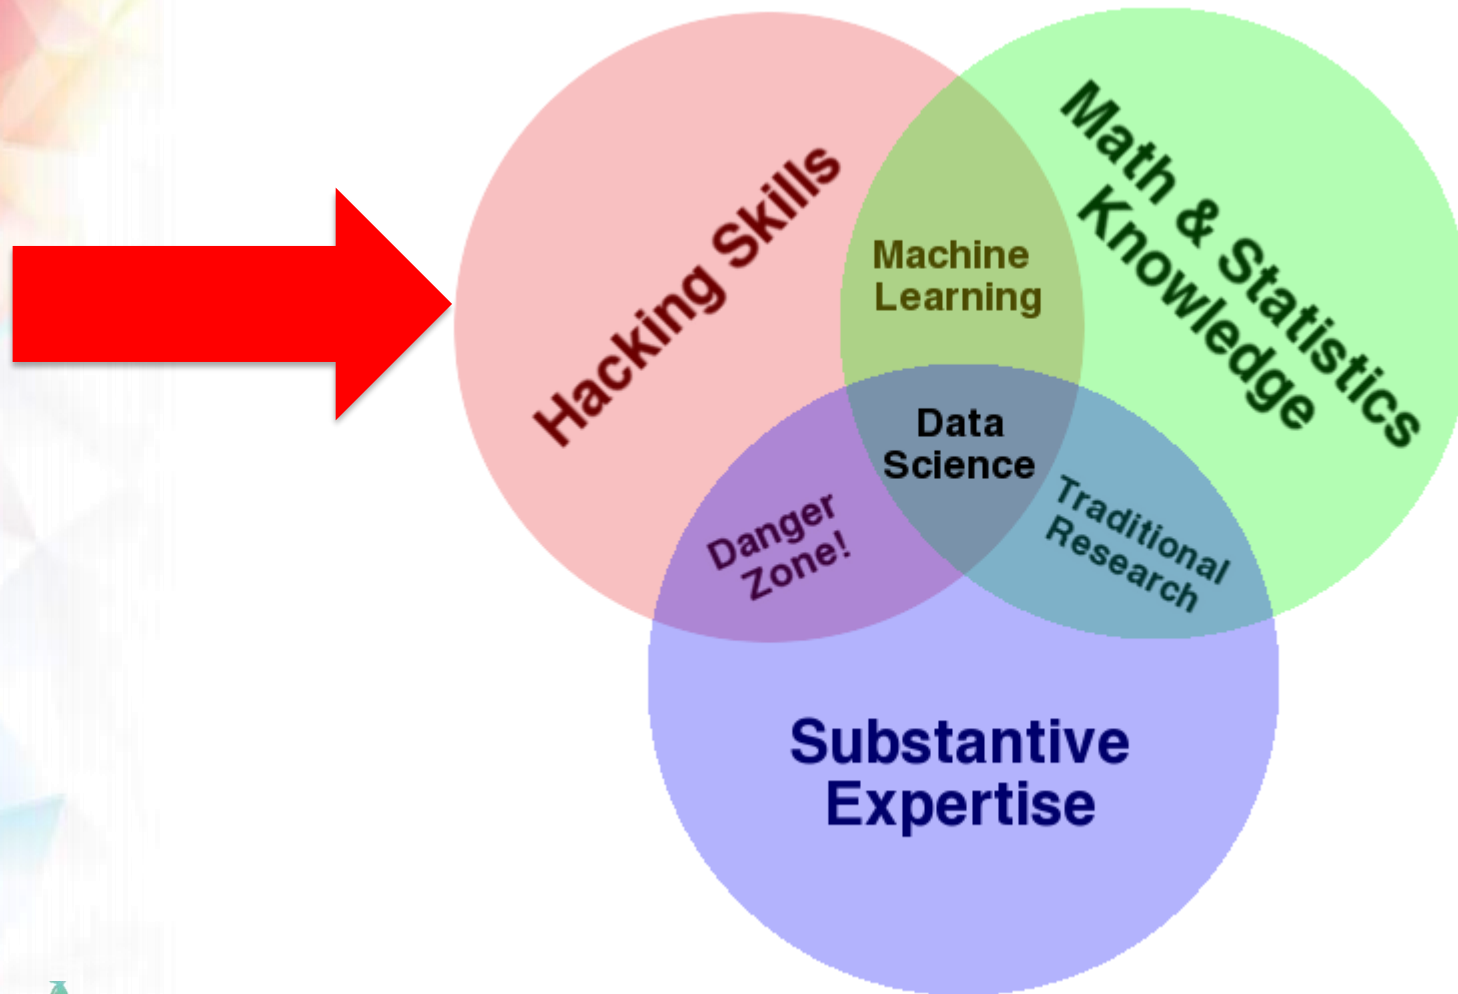
Analyzing the Analysts

		Hacker																Scripter					Application User												
		Analytics	Biotech	Datamart	Finance	Finance	Healthcare	Healthcare	Healthcare	Insurance	Marketing	Marketing	News	Retail	Retail	Social Aggregator	Social Networking	Social Networking	Visualization	Web	Web	Analytics	Analytics	Analytics	Finance	Healthcare	Media	Retail	Finance	Insurance	Retail	Retail	Sports	Web	Security
Process	Discovery	Locating Data	x	x	x	x	x	x	x	x																									
		Field Definitions	x	x	x	x	x	x	x																										
	Wrangle	Data Integration	x	x	x	x	x	x	x	x	x																								
		Parsing Semi-Structured	x	x	x																														
		Advanced Aggregation and Filtering	x																																
	Profile	Data Quality	x	x		x	x	x	x	x	x																								
		Verifying Assumptions		x																															
	Model	Feature Selection	x	x	x																														
		Scale	x	x	x	x	x																												
		Advanced Analytics	x	x																															
Report	Communicating Assumptions																																		
	Static Reports		x	x																															
Workflow	Data Migration	x	x	x																															
	Operationalizing Workflows	x	x																																
Tools	Database	SQL	x	x	x	x	x	x	x	x	x	x	x	x	x	x	x	x	x	x	x	x	x												
		Hadoop/Hive/Pig	x		x																														
		MongoDB																																	
		CustomDB	x																																
	Scripting	Java	x																																
		Perl																																	
		Python	x																																
		Clojure																																	
		Visual Basic																																	
	Modeling	R	x																																
		Matlab																																	
		SAS	x																																
		Excel																																	

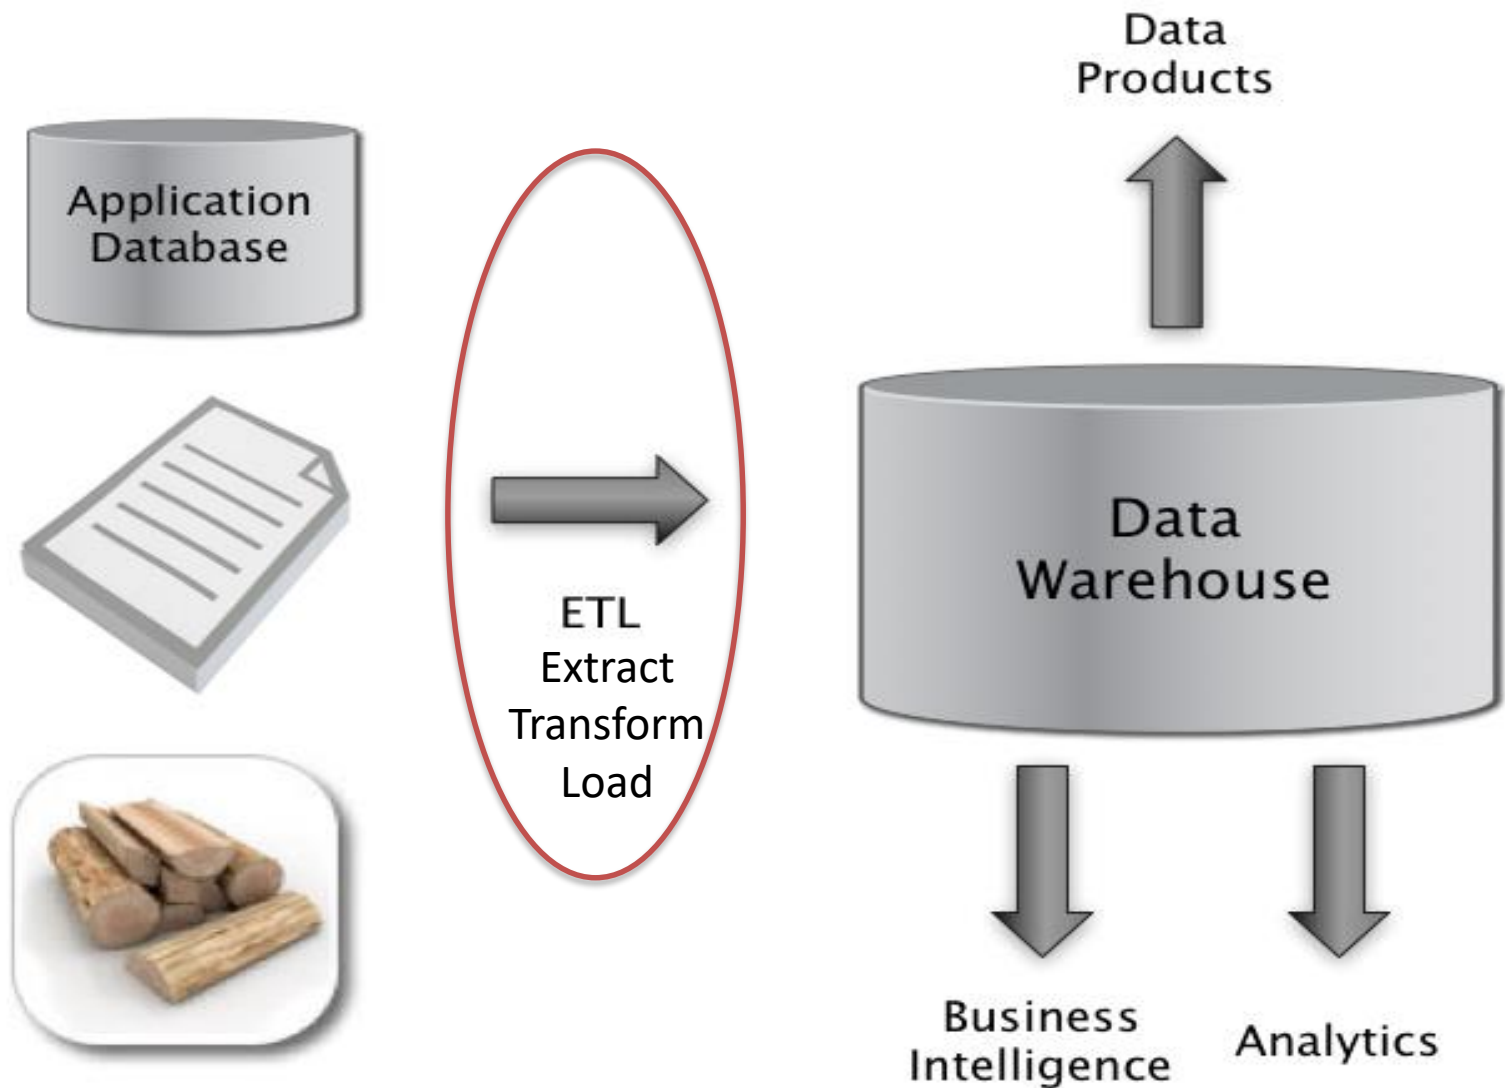
Fig. 1. Respondents, Challenges and Tools. The matrix displays interviewees (grouped by archetype and sector) and their corresponding challenges and tools. *Hackers* faced the most diverse set of challenges, corresponding to the diversity of their workflows and toolset. *Application users* and *scripters* typically relied on the IT team to perform certain tasks and therefore did not perceive them as challenges.

From Kandel, Paepcke, Hellerstein and Heer, "Enterprise Data Analysts and Visualization: An Interview Study", IEEE VAST 2012

Data Science – One Definition



The Big Picture



Data Preparation overview

- ETL
 - We need to **extract** data from the **source(s)**
 - We need to **load** data into the **sink**
 - We need to **transform** data at the source, sink, or in a **staging area**
 - Sources: file, database, event log, web site, HDFS...
 - Sinks: Python, R, SQLite, RDBMS, NoSQL store, files, HDFS...

Data Sources at Web Companies

- Examples from Facebook
 - Application databases
 - Web server logs
 - Event logs
 - API server logs
 - Ad server logs
 - Search server logs
 - Advertisement landing page content
 - Wikipedia
 - Images and video
- } Structured Data
- } Semi-structured Data
- } Unstructured Data

The (changing) role of Schema

Schema specify the **structure** and **types** of a data repository, e.g. the types of each column in a table.

They may also specify constraints **within** or **between** data fields.

Traditional databases are **schema-on-write**. You cannot load data into a table without a schema.

Newer (noSQL) data stores are **schema-on-read** or **schemaless**: You can defer applying a schema until you read the data, or avoid schema altogether.

Schema-on-Write

SQL:

```
CREATE SCHEMA Sprockets
```

```
CREATE TABLE NineProngs (source int, cost int, partnumber int)  
GO
```

```
INSERT INTO NineProngs (source, cost, partnumber)  
VALUES (5, 100, 45312453)
```

Schema-on-Read Data Types

XML: Generalizes HTML and specifies data **structure**. XML schema can be applied later to interpret XML data and specify **data types**. Here is some XML-encoded **data**:

```
<location>  
  <latitude>37.78333</latitude>  
  <longitude>122.4167</longitude>  
</location>
```

When stored without a schema, the numerical data are stored as **strings**.

XML Schema

```
<location>
  <latitude>37.78333</latitude>
  <longitude>122.4167</longitude>
</location>
```

An XML schema for this element:

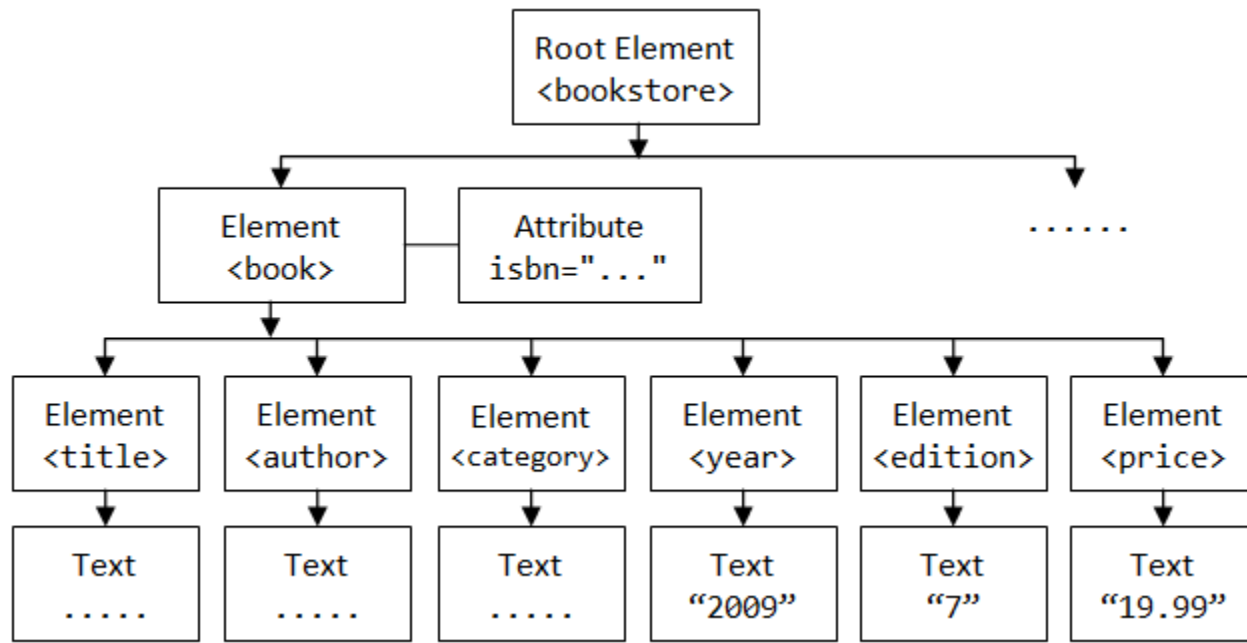
...

```
<xsd:schema xmlns:xsd="http://www.w3.org/2001/XMLSchema"
  elementFormDefault="unqualified">
  <xsd:complexType name="location">
    <xsd:sequence>
      <xsd:element name="latitude" type="xsd:decimal"/>
      <xsd:element name="longitude" type="xsd:decimal"/>
    </xsd:sequence>
  </xsd:complexType name="location">
```

XML and DOM

XML is a text format that encodes DOM (Document-Object Models) which is a data structure e.g. for Web pages.

The DOM is tree-structured:



XML Queries

XML schema allow a DB to interpret the data when running queries, e.g. to do **arithmetic** or **range queries** on numerical values.

XQuery is a standard for querying XML data with or without schema:

```
<places>{  
  for $city in /map/city  
    let $latlong := $city/location  
    where ((xs:float($latlong/latitude) < 39) and  
           (xs:float($latlong/latitude) > 38))  
    return  
      <place name="{ $city/name}"/>  
}</places>
```

JSON

JSON (Javascript Object Notation) by contrast is a schemaless data description language (Schema support was added later):

```
{
  "firstName": "John",
  "lastName": "Smith",
  "age": 25,
  "address": {
    "streetAddress": "21 2nd Street",
    "city": "New York",
    "state": "NY",
    "postalCode": "10021-3100" },
  "phoneNumbers": [
    { "type": "home",
      "number": "212 555-1234" },
    { "type": "office",
      "number": "646 555-4567" } ],
  "children": [],
  "spouse": null
}
```


JSON

JSON is typically used to represent hierarchical data structures directly in the target language (Javascript or Java).

Transformations on the data are **procedural** in the target language (not declarative in a language as in Xquery).

Easier for some tasks, but painful for e.g. schema changes.

Data Tools

XML:

- Separation between schema and data.
- Data can be represented and stored without schema (as strings).
- More verbose (but not true after compression or in DB).
- Standard Query/Transformation languages XSLT and Xquery.

JSON:

- Types inferred inline. Schema rarely used but can be.
- Data without schema use type inference (string, int, float,...).
- More succinct in ASCII form.
- Transformation/ingestion rely on code (Java or Javascript).

Data Tools

XML:



- Mark Logic Server
 - XQuery-based, semi-structured data, late/early Schema use
 - Also many traditional DB features: transactions, journaling, fine-grained access control,...

JSON:



- MongoDB
 - JSON native, “schemaless”
 - Based on Open-source code

Log Files – Example Apache Web Log

Processes, usually daemons, create logs
e.g., httpd, mysqld, syslogd

- 66.249.65.107 - - [08/Oct/2007:04:54:20 -0400] "GET /support.html HTTP/1.1" 200 11179 "-" "Mozilla/5.0 (compatible; Googlebot/2.1; +http://www.google.com/bot.html)"
- 111.111.111.111 - - [08/Oct/2007:11:17:55 -0400] "GET / HTTP/1.1" 200 10801 "http://www.google.com/search?q=log+analyzer&ie=utf-8&oe=utf-8 &aq=t&rls=org.mozilla:en-US:official&client=firefox-a" "Mozilla/5.0 (Windows; U; Windows NT 5.2; en-US; rv:1.8.1.7) Gecko/20070914 Firefox/2.0.0.7"
- 111.111.111.111 - - [08/Oct/2007:11:17:55 -0400] "GET /style.css HTTP/1.1" 200 3225 "\"<http://www.loganalyzer.net/>\" "Mozilla/5.0 (Windows; U; Windows NT 5.2; en-US; rv:1.8.1.7) Gecko/20070914 Firefox/2.0.0.7"

Tabular Data

- What is a table?
 - A **table** is a collection of **rows** and **columns**
 - Each row has an **index**
 - Each column has a **name**
 - A **cell** is specified by an (index, name) pair
 - A cell may or may not have a **value**
- Schema = (minimally) column types.
- Often stored as text files in CSV or TSV format.

Tabular Data

- Fortune 500

	A	B	C	D	E	F	G	H	I
1	rank	company	cik	ticker	sic	state_location	state_of_incorporation	revenues	profits
2	1	Wal-Mart Stores	104169	WMT	5331	AR	DE	421849	16389
3	2	Exxon Mobil	34088	XOM	2911	TX	NJ	354674	30460
4	3	Chevron	93410	CVX	2911	CA	DE	196337	19024
5	4	ConocoPhillips	1163165	COP	2911	TX	DE	184966	11358
6	5	Fannie Mae	310522	FNM	6111	DC	DC	153825	-14014
7	6	General Electric	40545	GE	3600	CT	NY	151628	11644
8	7	Berkshire Hathaway	1067983	BRKA	6331	NE	DE	136185	12967
9	8	General Motors	1467858	GM	3711	MI	MI	135592	6172
10	9	Bank of America Corp.	70858	BAC	6021	NC	DE	134194	-2238
11	10	Ford Motor	37996	F	3711	MI	DE	128954	6561
12	11	Hewlett-Packard	47217	HPQ	3570	CA	DE	126033	8761
13	12	AT&T	732717	T	4813	TX	DE	124629	19864
14	13	J.P. Morgan Chase & Co.	19617	JPM	6021	NY	DE	115475	17370
15	14	Citigroup	831001	C	6021	NY	DE	111055	10602
16	15	McKesson	927653	MCK	5122	CA	DE	108702	1263
17	16	Verizon Communications	732712	VZ	4813	NY	DE	106565	2549
18	17	American International Group	5272	AIG	6331	NY	DE	104417	7786
19	18	International Business Machines	51143	IBM	3570	NY	NY	99870	14833
20	19	Cardinal Health	721371	CAH	5122	OH	OH	98601.9	642.2
21	20	Freddie Mac	37785	FMC	2800	PA	DE	98368	-14025

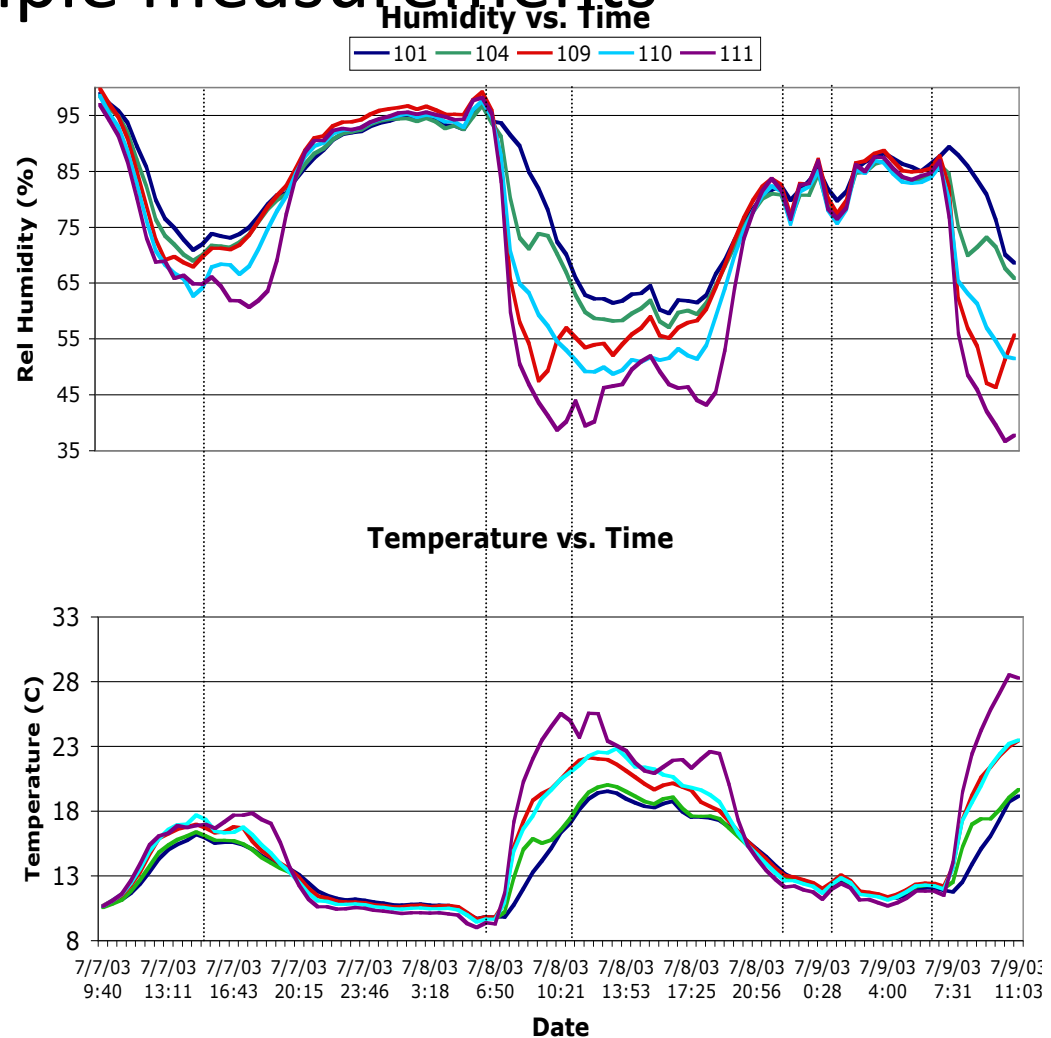

Tabular Data (csv)

- Fortune 500

Fortune 500 with ticker and EDGAR - Plus Ticker and EDGAR.txt

rank	company	cik	ticker	sic	state	location	state_of_incorporation	revenues	profits
1	Wal-Mart Stores	104169	WMT	5331	AR	DE	421849	16389	
2	Exxon Mobil	34088	XOM	2911	TX	NJ	354674	30460	
3	Chevron	93410	CVX	2911	CA	DE	196337	19024	
4	ConocoPhillips	1163165	COP	2911	TX	DE	184966	11358	
5	Fannie Mae	310522	FNM	6111	DC	DC	153825	-14014	
6	General Electric	40545	GE	3600	CT	NY	151628	11644	
7	Berkshire Hathaway	1067983	BRKA	6331	NE	DE	136185	12967	
8	General Motors	1467058	GM	3711	MI	MI	135592	6172	
9	Bank of America Corp.	70058	BAC	6021	NC	DE	134194	-2238	
10	Ford Motor	37996	F	3711	MI	DE	128954	6561	
11	Hewlett-Packard	47217	HPQ	3570	CA	DE	126033	8761	
12	AT&T	732717	T	4813	TX	DE	124629	19064	
13	J.P. Morgan Chase & Co.	19617	JPM	6021	NY	DE	115475	17370	
14	Citigroup	831001	C	6021	NY	DE	111055	10602	
15	McKesson	927653	MCK	5122	CA	DE	108702	1263	
16	Verizon Communications	732712	VZ	4813	NY	DE	106565	2549	
17	American International Group	5272	AIG	6331	NY	DE	104417	7786	
18	International Business Machines	51143	IBM	3570	NY	NY	99870	14833	
19	Cardinal Health	721371	CAH	5122	OH	OH	98601.9	642.2	
20	Freddie Mac	37785	FMC	2800	PA	DE	98368	-14025	
21	CVS Caremark	64803	CVS	5912	RI	DE	96413	3427	
22	UnitedHealth Group	731766	UNH	6324	MN	MN	94155	4634	
23	Wells Fargo	72971	WFC	6021	CA	DE	93249	12362	
24	Valero Energy	1035002	VLO	2911	TX	DE	86034	324	
25	Kroger	56873	KR	5411	OH	OH	82189.4	1116.3	
26	Procter & Gamble	80424	PG	2840	OH	OH	79689	12736	
27	AmerisourceBergen	1140859	ABC	5122	PA	DE	77954	636.7	
28	Costco Wholesale	909832	COST	5331	WA	WA	77946	1303	
29	Marathon Oil	101778	MRO	2911	TX	DE	68413	2568	
30	Home Depot	354950	HD	5211	GA	DE	67997	3338	

Internet of Things: Example measurements



Tabular Data from Sensors

Goals

- Want to support both long-term (trend) and short-term (real-time) queries.
- Want low latency but also efficient, real-time indexing for longer-term queries.
- Want triggers (alerts) for a variety of conditions.

Tabular Data from Sensors

Tools:

- Microsoft SQL server, Oracle

Analysis:

- Matlab still widely used for analysis in Financial services, Python tools.

Syslog – A Standard for System Messages

- Developed by Eric Allman (at Berkeley) as part of the Sendmail project
- Standardized by the IETF in RFC 3164 and RFC 5424
- Listens on port 514 using UDP
- Puts data in /var/log/messages by default
- Enables rich analysis:

splunk>



Syslog

dhcp-47-129:DataScienceF15> syslog -w 10

Feb 3 15:18:11 dhcp-47-129 Evernote[1140] <Warning>: -[EDAMAccounting read:]: unexpected field ID 23 with type 8. Skipping.

Feb 3 15:18:11 dhcp-47-129 Evernote[1140] <Warning>: -[EDAMUser read:]: unexpected field ID 17 with type 12. Skipping.

Feb 3 15:18:11 dhcp-47-129 Evernote[1140] <Warning>: -[EDAMAuthenticationResult read:]: unexpected field ID 6 with type 11. Skipping.

Feb 3 15:18:11 dhcp-47-129 Evernote[1140] <Warning>: -[EDAMAuthenticationResult read:]: unexpected field ID 7 with type 11. Skipping.

Feb 3 15:18:11 dhcp-47-129 Evernote[1140] <Warning>: -[EDAMAccounting read:]: unexpected field ID 19 with type 8. Skipping.

Feb 3 15:18:11 dhcp-47-129 Evernote[1140] <Warning>: -[EDAMAccounting read:]: unexpected field ID 23 with type 8. Skipping.

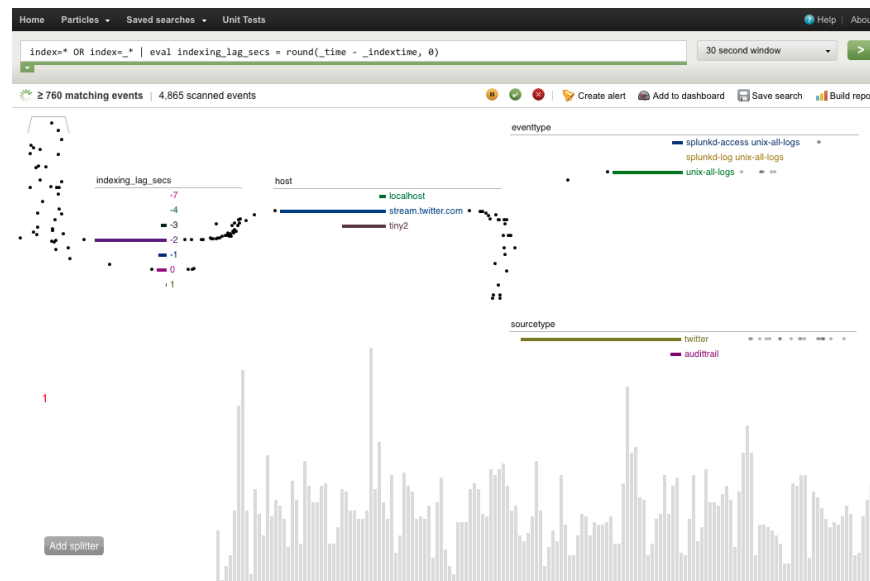
Feb 3 15:18:11 dhcp-47-129 Evernote[1140] <Warning>: -[EDAMUser read:]: unexpected field ID 17 with type 12. Skipping.

Feb 3 15:18:11 dhcp-47-129 Evernote[1140] <Warning>: -[EDAMSyncState read:]: unexpected field ID 5 with type 10. Skipping.

Feb 3 15:18:49 dhcp-47-129 com.apple.mtmd[47] <Notice>: low priority thinning needed for volume Macintosh HD (/) with 18.9 <= 20.0 pct free space

“Splunking”

- Grab data from many machines
- Index it
- Check for unusual events:
 - Disk problems
 - Network congestion
 - Security attacks
- Monitor Resources
 - Network
 - Memory usage
 - Disk use, latency
 - Threads
- Dashboard for cloud administration.



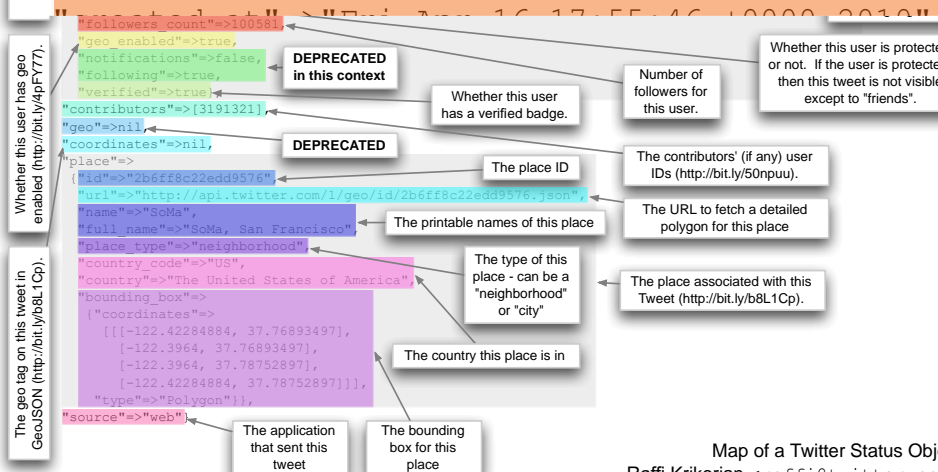
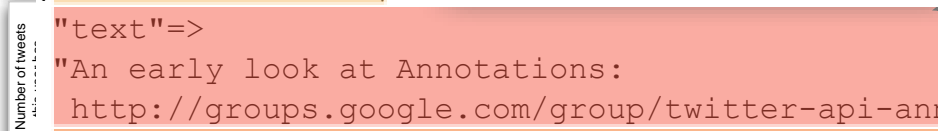

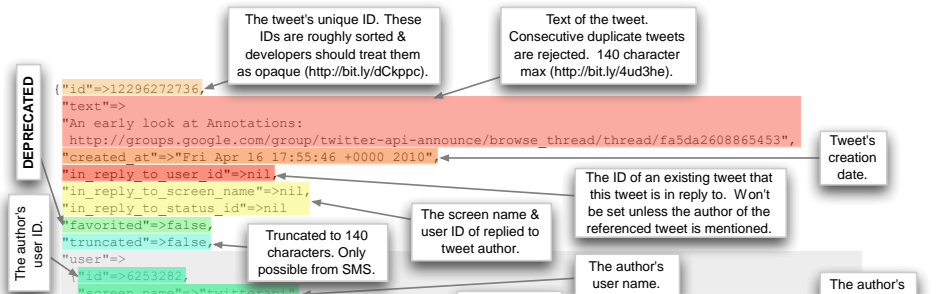
Some Questions

1) How Many Characters are there in a Tweet?

2) How Many Bytes are there in the API record for a Tweet?



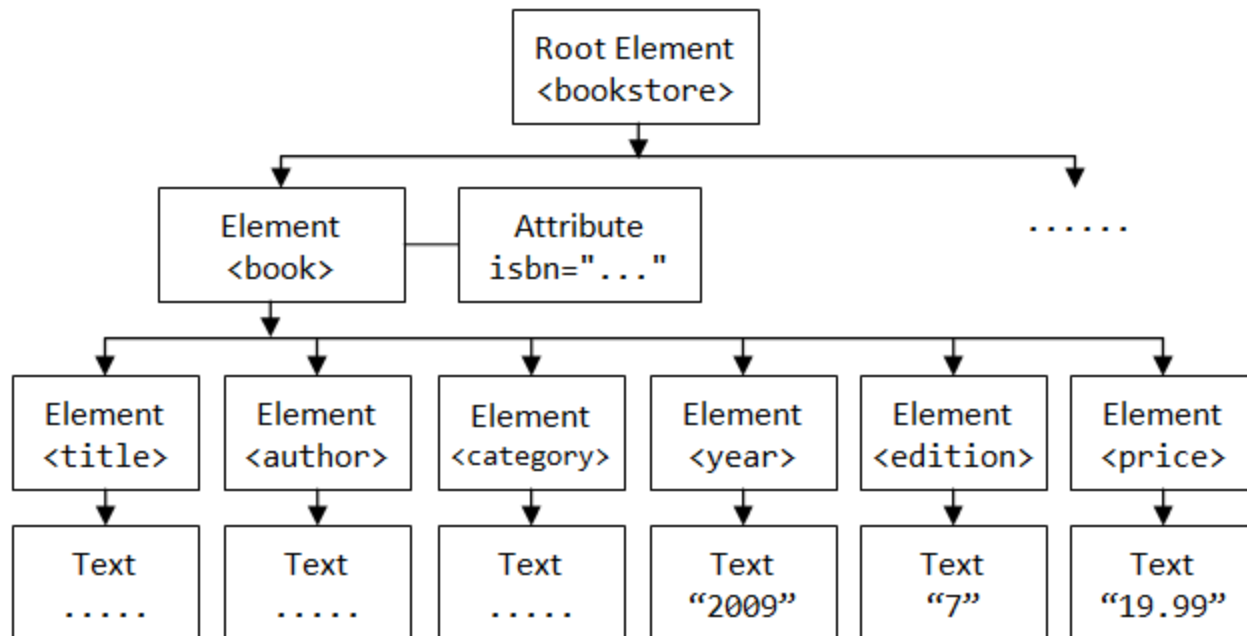
Tweet JSON Format



Text of the tweet.
Consecutive duplicate tweets
are rejected. 140 character
max (<http://bit.ly/4ud3he>).

Processing XML and JSON

- The DOM is an easy object to work with: all the data in the object is accessible by links.
- The problem is that I might not care about most of the data, and I might not be able to fit the DOM for a large object in RAM.



Event-Driven Parsing: SAX

<?xml version="1.0" encoding="UTF-8"?> → Document Header

<!-- bookstore.xml --> → Comment

<bookstore> → Start-element "bookstore"

<book ISBN="0123456001"> → Start-element "book"

<title>Java For Dummies</title> → Start-element "title"

<author>Tan Ah Teck</author> End-element "title"

<category>Programming</category>

<year>2009</year>

<edition>7</edition>

<price>19.99</price>

</book> → End-element "book"

Event-Driven Parsing: SAX

A SAX parser finds all the open-close-tag events in an XML documents, and does callbacks to user code.

- User code can respond to only a subset of events corresponding to the tags it is interested in.
- User code can correctly compute aggregates from the data rather than create a record for each tag.
- User code must implement a state machine to keep track of “where it is” in the DOM tree.
- User code can implement flexible error recover strategies for ill-formed XML.



What about JSON?

Most JSON parsers construct the “DOM” directly.

But there are a few SAX-style parsers:

- Jackson
- JSON-simple

What about HTML?

- Common Crawl, **about 5 billion web pages**, between **0.2-0.5%** of Google's web crawl.
- 60 TB, hosted on Amazon S3, also available for download.
- Includes **link data, page rank**.
- In ARC (Internet Archive) File format.
- So there's plenty of data, and there are many crawlers for targeted exploration...
 - HTTrack, ...



HTML Tag Soup

```
<!DOCTYPE html PUBLIC "-//W3C//DTD XHTML 1.0 Transitional//EN"
"http://www.w3.org/TR/xhtml1/DTD/xhtml1-transitional.dtd">
<html xmlns="http://www.w3.org/1999/xhtml">
<head><!-- types/widgets/pages/common/page.tpl home/index_v3.html generated by index_v3 on Wed 29
Feb 2012 11:04:41 PM PST -->
<title>San Francisco Bay Area &mdash; News, Sports, Business, Entertainment, Classifieds: SFGate</title>
<meta http-equiv="content-type" content="text/html; charset=iso-8859-1" />
<meta name="description" content="Find local news &amp; information, updated weather, traffic, classifieds,
sports scores, real estate, jobs, cars, food &amp; wine, travel, entertainment, events and more on SFGate.com.
Connect to the Bay Area community." />
<meta name="keywords" content="San Francisco, San Francisco Bay Area, news, local events, breaking news,
world news, San Francisco Chronicle, SFGate" />
<meta property="fb:page_id" content="105702905593" />
<meta property="fb:admins" content="653226748,658759748" />
<!-- /widgets/sitewide/css/all/inc.html widgets/pages/common/post_write_mtime/css_inc.tpl -->
<!-- generated by sitewidecss on Thu 16 Feb 2012 10:41:53 AM PST -->
<link rel="stylesheet" type="text/css" title="SFGate" media="all"
href="http://imgs.sfgate.com/css1329417713/sitewide/css/sitewide.css" />
<!-- sitewide/css/all/inc.html end css_inc.tpl -->
```

HTML Tools - Parsing

- “Beautiful Soup”
<http://www.crummy.com/software/BeautifulSoup/>
a Python API for handling real HTML. DOM or SAX interfaces.
- “TagSoup”
<http://ccil.org/~cowan/XML/tagsoup/>
provides a Sax interface, i.e. a streaming parse, to Java applications. Can transform to a format you want using XSLT.
- Taggle, part of the Arabica toolset
<http://www.jezuk.co.uk/cgi-bin/view/arabica/code>
is a version of TagSoup written in C++. You may want to use this if you have a lot of data.



Web Services

Most large web sites today actively discourage screen-scraping to get their content, and provide Web Service APIs instead.

This is the “right” way to get data from online sources.



Web Services

W3C definition:

a "Web service" as "a software system designed to support interoperable machine-to-machine interaction over a network".

Two kinds:

- XML-based RPC-style messages: SOAP
- REST-style stateless interactions, URLs encode state

Can run over different transports, but usually HTTP

Examples

Twitter: REST API and streaming API with JSON content. Provides sampling, searching and filtering capabilities.

Amazon: has a “product advertising API” in XML with a WSDL spec. Includes product search, reviews etc.

Livejournal: RSS/Atom + custom XML/RPC. Search by keyword, topic, follow friend links.

Netflix: Javascript, Atom and REST interfaces.

Ebay: Many APIs for searching, buying and posting. WSDL descriptions, client code in Java and .NET

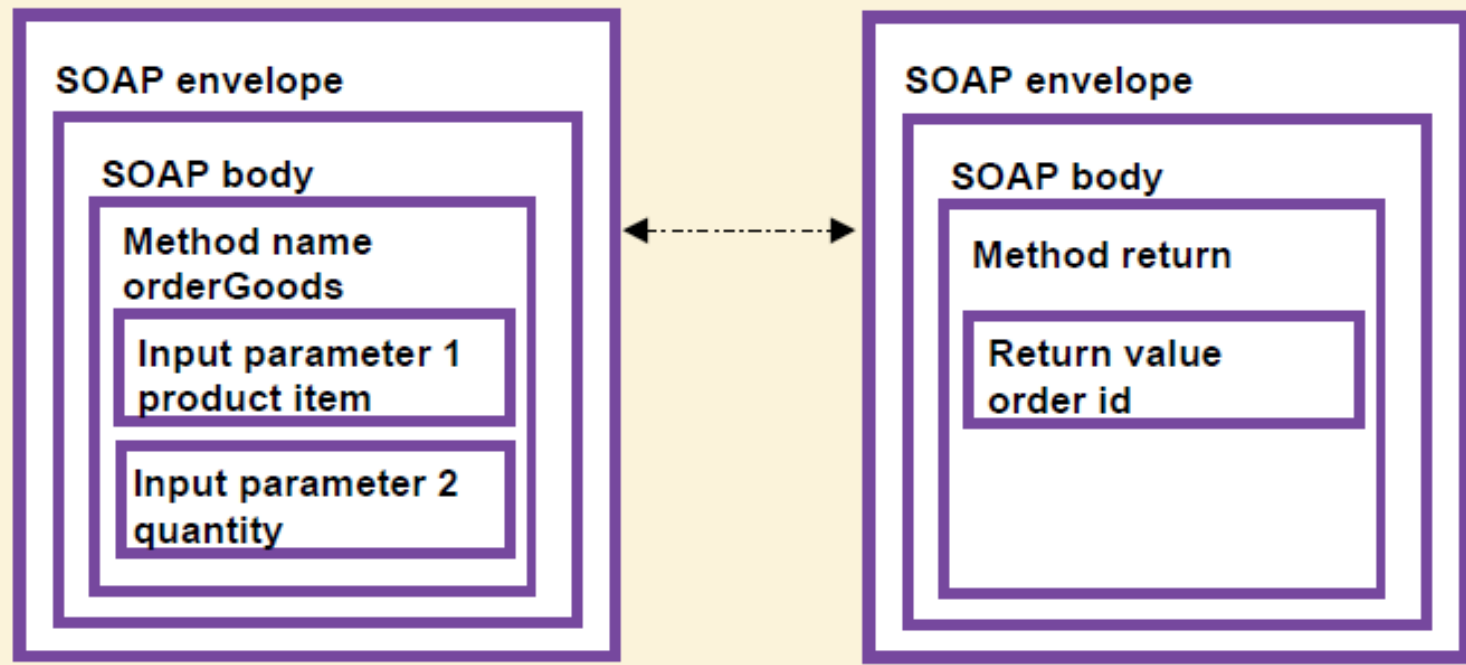
Flickr: Comprehensive API set, free for non-commercial use. REST, XML-RPC, SOAP, with client code in many languages.

vBulletin: REST interface, most actions supported



SOAP RPC

SOAP RPC messages typically encode arguments that are presented to the calling program as parameters and return values. HTTP POST/GET are used to communicate:



Soap RPC

```
POST /travelservice
SOAPAction: "http://www.acme-travel.com/flightinfo"
Content-Type: text/xml; charset="utf-8"
Content-Length: nnnn

<SOAP:Envelope xmlns:SOAP=
  "http://schemas.xmlsoap.org/soap/envelope/">
  <SOAP:Body>
    <m:GetFlightInfo
      xmlns:m="http://www.acme-travel.com/flightinfo"
      SOAP:encodingStyle=
        "http://schemas.xmlsoap.org/soap/encoding/"
      xmlns:xsd="http://www.w3.org/2001/XMLSchema"
      xmlns:xsi=
        "http://www.w3.org/2001/XMLSchema-instance">
      <airlineName xsi:type="xsd:string">UL
    </airlineName>
      <flightNumber xsi:type="xsd:int">506
    </flightNumber>
    </m:GetFlightInfo>
  </SOAP:Body>
</SOAP:Envelope>
```

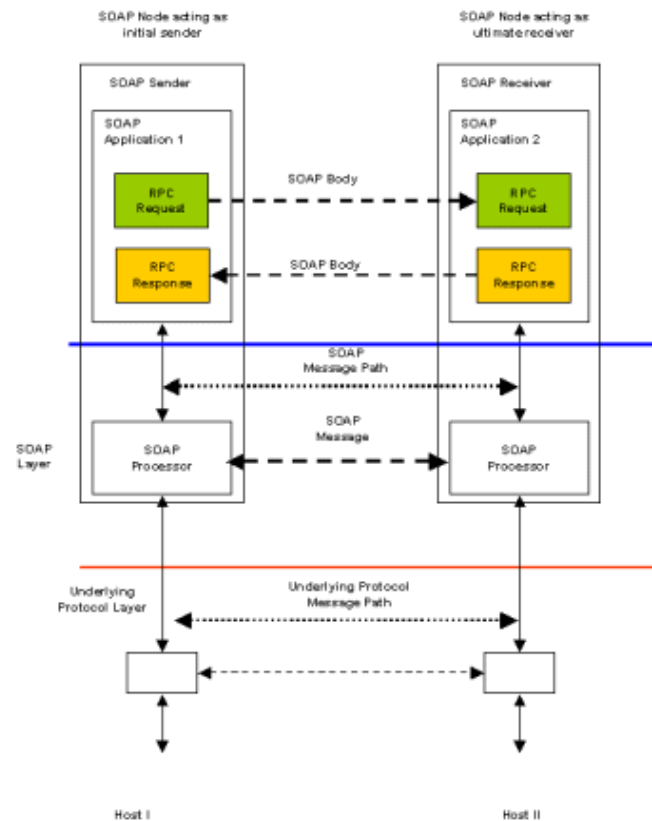
Soap Response

```
HTTP/1.1 200 OK
Content-Type: text/xml; charset="utf-8"
Content-Length: nnnn

<SOAP:Envelope xmlns:SOAP=
  "http://schemas.xmlsoap.org/soap/envelope/">
  <SOAP:Body>
    <m:GetFlightInfoResponse
      xmlns:m="http://www.acme-travel.com/flightinfo"
      SOAP:encodingStyle=
        "http://schemas.xmlsoap.org/soap/encoding/"
      xmlns:xsd="http://www.w3.org/2001/XMLSchema"
      xmlns:xsi=
        "http://www.w3.org/2001/XMLSchema-instance">
      <flightInfo>
        <gate xsi:type="xsd:int">10</gate>
        <status xsi:type="xsd:string">ON TIME</status>
      </flightInfo>
    </m:GetFlightInfoResponse>
  </SOAP:Body>
</SOAP:Envelope>
```

Web Services

XML-RPC, requires a request-response cycle. Often longer “conversations.” i.e. it’s a stateful protocol, and both endpoints need to agree on the state.



REST

REpresentation State Transfer

Stateless Client/Server Protocol: Principles

1. Each message in the protocol contains all the information needed by the receiver to understand and/or process it. This constraint attempts to “**keep things simple**” and avoid needless complexity
2. Set of Uniquely Addressable Resources
 - “**Everything is a Resource**” in a RESTful system
 - Requires universal syntax for resource identification (e.g. URI)

REST

3. Set of Well-Defined Operations that can be applied to all resources

- In context of HTTP, the primary methods are
- POST, GET, PUT, DELETE
- these are similar (but not exactly) to the database notion of
- CRUD (Create, Read, Update, Delete)

4. The use of Hypermedia both for Application Information and State Transitions

- Resources are typically stored in a structured data format that supports hypermedia links, such as XHTML or XML



REST example

```
<user>
```

```
  <name>Jane</name>
```

```
  <gender>female</gender>
```

```
  <location href="http://www.example.org/us/ny/new_york">
```

```
    New York City, NY, USA</location>
```

```
</user>
```

This documentation is a representation used for the User resource

It might live at <http://www.example.org/users/jane/>

- If a user needs information about Jane, they GET this resource
- If they need to modify it, they GET it, modify it, and PUT it back
- The href to the Location resource allows savvy clients to gain access to its information with another simple GET request

Implication: Clients cannot be too “thin”; need to understand resource formats



REST vs. RPC

In RPC systems, the design emphasis is on **verbs**

- What operations can I invoke on a system?
- getUser(), addUser(), removeUser(), updateUser(), getLocation(), updateLocation(), listUsers(), listLocations(), etc.

In REST systems, the design emphasis is on **nouns**

- User, Location
- In REST, you would define XML representations for these resources and then apply the standard methods to them



Dirty Data


- The **Statistics** View:
 - There is a process that produces data
 - We want to model ideal samples of that process, but in practice we have non-ideal samples:
 - **Distortion** – some samples are corrupted by a process
 - **Selection Bias** - likelihood of a sample depends on its value
 - **Left and right censorship** - users come and go from our scrutiny
 - **Dependence** – samples are supposed to be independent, but are not (e.g. social networks)
 - You can add new models for each type of imperfection, but you can't model everything.
 - What's the best trade-off between accuracy and simplicity?



Numeric Outliers

12	13	14	21	22	26	33	35	36	37	39	42	45	47	54	57	61	68	450
----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	-----

ages of employees (US)



median 37

mean 58.52632

variance 9252.041

Adapted from Joe Hellerstein's 2012 CS 194 Guest Lecture

Data Cleaning Tools: OpenRefine

- Spreadsheet-like tool allowing data quality checking: reformatting, substitution, constraint checking etc.

11285 rows

Extensions:

Zemanta

Freebase

RDF

CK

Show as:

rows

records

 Show:

5

10

25

50

 rows

« first < previous 1 - 50 next >

<div><input type="checkbox"/> All</div>	<div><input type="checkbox"/> Capital or Revenue</div>	<div><input type="checkbox"/> Directorate</div>	<div><input type="checkbox"/> Transaction Number</div>	<div><input type="checkbox"/> Date</div>	<div><input type="checkbox"/> Service Area</div>	<div><input type="checkbox"/> Expenses Type</div>	<div><input type="checkbox"/> Amount</div>	<div><input type="checkbox"/> Supplier</div>
<div>☆</div> <div>🔊</div> 1. Revenue Community Wellbeing & Social Care 5105695746 05.04.2013 Youth & Community Operational Equipment 120 REDACTED PERSON/								
<div>☆</div> <div>🔊</div> 2. Revenue Community Wellbeing & Social Care 5105695746 05.04.2013 Youth & Community Operational Equipment 80 REDACTED PERSON/								
<div>☆</div> <div>🔊</div> 3. Revenue Community Wellbeing & Social Care 5105698650 24.04.2013 Leaseholds by LA Accommodation Costs - Leaseholder Payments edit 695.89 REDACTED PERSON/								
<div>☆</div> <div>🔊</div> 4. Revenue Community Wellbeing & Social Care 5105698650 24.04.2013 Leaseholds by LA Accommodation Costs - Leaseholder Payments 695.89 REDACTED PERSON/								
<div>☆</div> <div>🔊</div> 5. Revenue Community Wellbeing & Social Care 5105698650 24.04.2013 Leaseholds by LA Accommodation Costs - Leaseholder Payments 695.89 REDACTED PERSON/								
<div>☆</div> <div>🔊</div> 6. Revenue Community Wellbeing & Social Care 5105698650 24.04.2013 Leaseholds by LA Accommodation Costs - Leaseholder Payments 695.89 REDACTED PERSON/								
<div>☆</div> <div>🔊</div> 7. Revenue Community Wellbeing & Social Care 5105698650 24.04.2013 Leaseholds by LA Accommodation Costs - Leaseholder Payments 695.89 REDACTED PERSON/								
<div>☆</div> <div>🔊</div> 8. Revenue Chief Executive, Schools & Learning 5105698316 19.04.2013 L&A Commissioned Activity Bought in Prof Services - Curriculum (Schools) 250 REDACTED PERSON/								
<div>☆</div> <div>🔊</div> 9. Revenue Chief Executive, Schools & Learning 5105698318 19.04.2013 L&A Commissioned Activity Bought in Prof Services - Curriculum (Schools) 710 REDACTED PERSON/								
<div>☆</div> <div>🔊</div> 10. Revenue Economy & Environment 5105695879 05.04.2013 IW Biological Record General Materials 220.2 REDACTED PERSON/								
<div>☆</div> <div>🔊</div> 11. Revenue Chief Executive, Schools & Learning 5105696514 12.04.2013 Adult Services Training Training and Conferences 150 REDACTED PERSON/								
<div>☆</div> <div>🔊</div> 12. Revenue Community Wellbeing & Social Care 5105695832 10.04.2013 Short Breaks Payments to Voluntary and Other Associations 1,260.00 REDACTED PERSON/								
<div>☆</div> <div>🔊</div> 13. Capital Resources 5105696504 12.04.2013 Capital Receipts External Design and Supervision Fees 400 REDACTED PERSON/								
<div>☆</div> <div>🔊</div> 14. Capital Resources 5105696505 12.04.2013 Capital Receipts External Design and Supervision Fees 1,350.00 REDACTED PERSON/								
<div>☆</div> <div>🔊</div> 15. Revenue Economy & Environment 5105696707 12.04.2013 Schools Reorganisation Security of Buildings 300 REDACTED PERSON/								
<div>☆</div> <div>🔊</div> 16. Revenue Economy & Environment 5105696707 12.04.2013 Schools Reorganisation Security of Buildings 300 REDACTED PERSON/								


Exploring

- Get familiar with your favorite graphing package:
 - Matplotlib is widely used in Python
 - Ggplot is good for more advanced plots (similar to R)
 - D3.js popular for interactive graphics, but low-level:
 - Bokeh provides high-level primitives
 - Vega/Vincent same goals, developed by Trifacta
- Get fluent with plotting:
 - Histograms
 - Scatter plots
 - Line and bar plots



Looking at Data


- Histograms can tell you a lot about a single variable, discrete or continuous:




Looking at Data

- Skewed distributions:

Skewed-Right Distribution




Skewed-Left Distribution




Long-tailed data

- Long tailed data



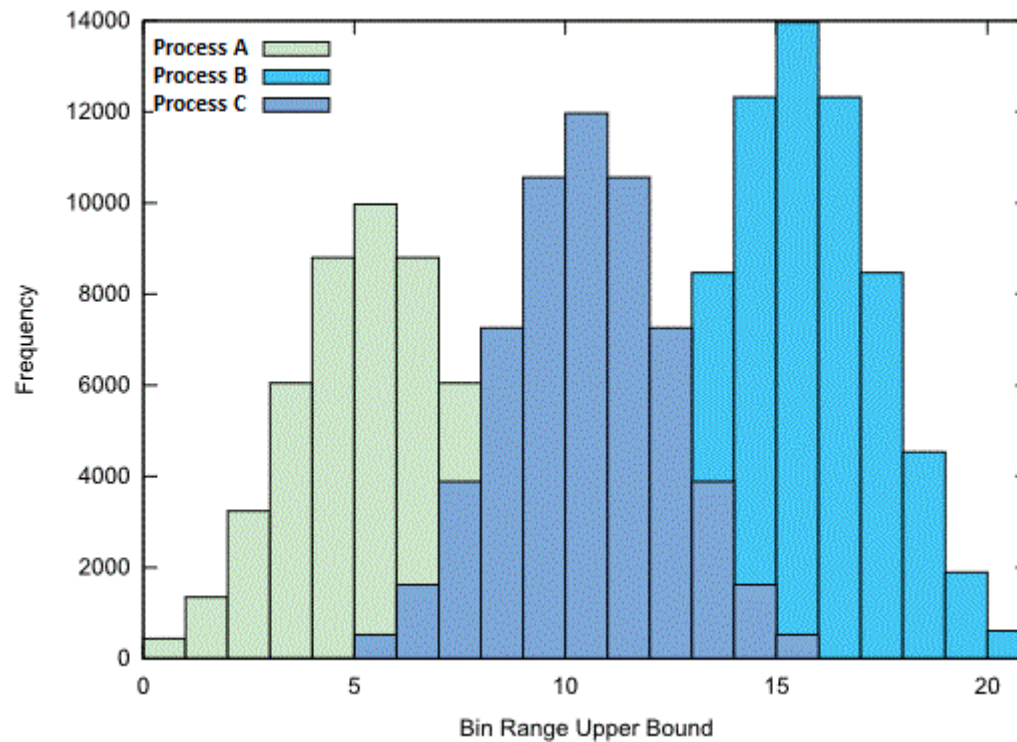
Multimodal data

- Two or more distinct peaks in a histogram.
- Suggests two or more distinct populations of samples.
- Often arise from gender/political views, other binary factors.
- But don't guess!! Explore further by using, e.g. color and a histogram of multiple populations.



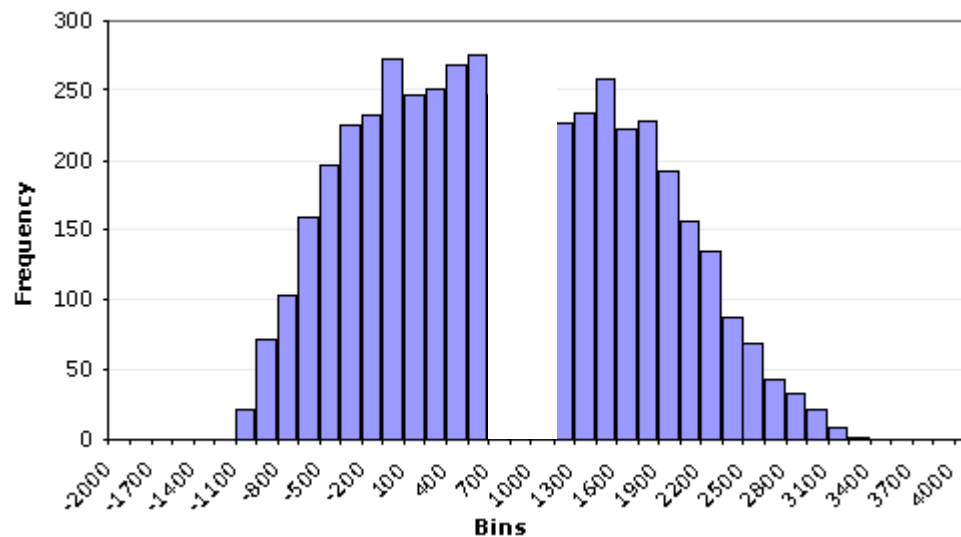
Multimodal data

- Explore further by using, e.g. color and a histogram of multiple populations.




Weird data

- Some data are very hard to explain.
- Don't try. Trace through the data pipeline to find where the strangeness comes from. Usually it's a processing bug.




Weird data

- Some data are very hard to explain.
- Don't try. Trace through the data pipeline to find where the strangeness comes from. Usually it's a processing bug.




Weird data

- Some data are very hard to explain.
- Don't try. Trace through the data pipeline to find where the strangeness comes from. Usually it's a processing bug.




Proactive Weird data Detection

- If data look normal, take a picture and save it for later...
- Then periodically compare new data with old whenever there is a pipeline update.
- Always try to have a theory of what the data should look like.





Two variables – Scatter plots

- Scatter plots quickly expose the relationships between two variables




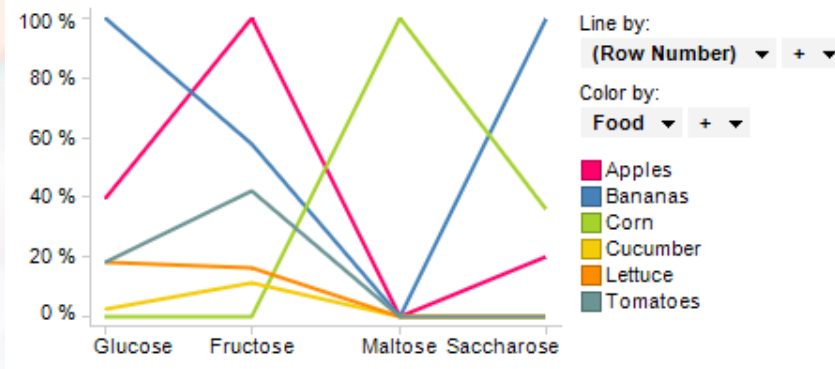
More than two variables

- **Stacked plot:** stack variable is discrete:




More than two variables

- **Parallel coordinate plot:** one discrete variable, an arbitrary number of other variables:



More than two variables

- **Radar Chart:** Similar: one discrete variable (design here), an arbitrary number of other variables:



Closing Remarks

- We argued for analysts to **form expectations** of what the data should look like. This helps guard against pipeline errors and to identify interesting patterns.
- An observer should also be atune to patterns that we not part of their theory. In other words to “expect the unexpected”.

